

Harmain Ali

New York, NY | harmain.dev@gmail.com | +1 (650) 555-1234 | <https://www.linkedin.com/in/harmainali-swe>

Professional Summary

Highly accomplished and results-driven Senior Software Engineer with 8+ years of experience in designing, developing, and deploying scalable, AI-powered distributed systems and complex infrastructure. Proven expertise in backend engineering, leveraging languages like Go and Python, and orchestrating robust solutions across cloud platforms including AWS, Azure, and GCP. Adept at implementing comprehensive DevOps practices, including CI/CD pipelines, containerization with Docker and Kubernetes, and Infrastructure as Code using Terraform and Ansible. Committed to building high-performance, observable systems, and developing innovative GTM tooling that drives business intelligence and operational efficiency.

Skills

Core Skills: Go, Python, C++, JavaScript, SQL, AWS, Azure, GCP, Kubernetes, Docker, Terraform, Ansible, CI/CD, Git, DevOps, Observability, Monitoring, Artificial Intelligence (AI), Machine Learning (ML), Distributed Systems, High-Availability Systems, Data Science, PyTorch, TensorFlow, Scikit-learn, NumPy, Pandas, ETL for LLMs, Vector Databases, Feature Stores, PostgreSQL, MongoDB (noSQL), SQL, RESTful APIs, React, Angular, Flask, ASP.Net Core, Jupyter Notebooks, VS Code, Google Colab, Linear Algebra, Algorithms, Statistics, Quality Control, Model Deployment, Data Analytics Automation

Professional Experience

AMD | New York, NY Jan 2024 - Dec 2026

Senior Software Engineer, AI Infrastructure

Technologies: Go, Python, Kubernetes, Terraform, Ansible, AWS, Azure, Docker, SQL, CI/CD, Observability

- Led the architectural design and implementation of scalable AI infrastructure, leveraging Kubernetes for container orchestration and Terraform for Infrastructure as Code across hybrid cloud environments including AWS and Azure, enhancing deployment efficiency by 30%.
- Developed critical backend services in Go and Python for real-time AI model inference and data processing, integrating with distributed systems to ensure high availability and low-latency performance.
- Engineered and optimized CI/CD pipelines using Ansible for automated provisioning and configuration management, significantly reducing deployment times and improving system reliability for AI-powered GTM tooling.
- Implemented advanced observability and monitoring solutions, including custom dashboards and alerting, to ensure production-grade availability and performance of AI systems, proactively identifying and resolving complex technical issues.

NVIDIA | New York, NY Jan 2021 - Dec 2023

Lead Backend Engineer, ML Platforms

Technologies: Go, Python, TensorFlow, Kubernetes, Terraform, GCP, AWS, React, PostgreSQL, RESTful APIs

- Designed and developed robust backend services in Go and Python to support large-scale machine learning platforms, integrating with GCP services for data storage and processing, improving system throughput by 25%.
- Orchestrated containerized ML workflows using Kubernetes, managing resource allocation and scaling for critical AI model training and deployment pipelines, reducing operational overhead by 20%.
- Implemented Infrastructure as Code (IaC) practices with Terraform to manage cloud resources on AWS and GCP, ensuring consistent and reproducible environments for ML development and production.

- Collaborated with front-end teams, providing RESTful APIs for React-based GTM tooling and internal applications, enhancing user interaction and data visualization for ML insights.

Google | New York, NY Jan 2018 - Dec 2020

Software Engineer, AI & Cloud Services

Technologies: Python, C++, Docker, GCP, SQL, Distributed Systems, Monitoring, ETL

- Developed and optimized core backend components for Google Cloud AI services using Python and C++, focusing on distributed systems and high-availability architectures, improving service uptime by 99.9%.
- Contributed to the design and implementation of automated data pipelines for AI systems, including ETL processes for large language models (LLMs) and feature stores, utilizing GCP services for scalable data management.
- Utilized Docker for containerization of AI applications, streamlining development and deployment processes across various environments and preparing for Kubernetes adoption.
- Implemented comprehensive monitoring and logging solutions for AI services, ensuring data quality and system performance in production, and providing insights for continuous improvement.

Certifications

- Certified Kubernetes Administrator (CKA)
- HashiCorp Certified: Terraform Associate
- Microsoft Certified: Azure Developer Associate
- Google Cloud Professional Cloud Architect
- Go Programming Language Certification
- Machine Learning Engineering with TensorFlow on Google Cloud | Google Cloud
- Deep Learning Specialization | DeepLearning.AI (Andrew Ng)

Education

Stanford University

Master of Science in Computer Science, specializing in Distributed Systems and Artificial Intelligence 2017

University of California, Berkeley

Bachelor of Science in Software Engineering 2016