

Dear Hiring Manager at datadog,

I am a seasoned backend engineer with six years of experience architecting and deploying scalable systems that process and analyze complex data at enterprise scale. As a Senior Machine Learning Engineer at AMD, I led the development of advanced AI models that improved processing efficiency by 15% while implementing automated monitoring systems that reduced critical errors by 20%—directly demonstrating my ability to build robust, reliable systems that handle mission-critical workloads. My background spans designing data pipelines, establishing quality control protocols, and collaborating across cross-functional teams to identify and solve complex engineering challenges, which aligns perfectly with Product Analytics' need to ingest, manipulate, correlate, and efficiently query data from multiple sources. Throughout my tenure at NVIDIA and Google, I developed and deployed production ML systems using Python, C++, and SQL, optimized data processing speeds by 25%, and architected scalable tool development pipelines that integrated complex algorithms—experiences that directly translate to designing and maintaining the core backend systems you're building. I am a hands-on technical leader comfortable influencing architecture decisions without direct authority, and I'm genuinely excited about leveraging AI tools to enhance problem-solving and code quality while contributing to building AI-enabled products. I'm particularly drawn to Datadog's hybrid culture and the opportunity to work on an early-stage, fast-growing product area where I can drive technical vision, mentor engineers, and deliver high-quality software that directly impacts how developers and product managers understand their users. I would welcome the opportunity to discuss how my experience building distributed systems, establishing data quality standards, and collaborating with platform teams can accelerate Product Analytics' technical roadmap.

Sincerely,

Harmain Ali