

Sabih Haider

San Francisco, CA | sabih0364@gmail.com | +923410084746 | <https://linkedin.com/in/sabih-haider>

Professional Summary

Full-stack software engineer with 8+ years of hands-on experience designing and building distributed systems, AI-powered infrastructure, and scalable cloud platforms. Proven expertise in building observability and monitoring solutions for complex systems, implementing infrastructure as code with Terraform and Ansible, and developing AI-augmented workflows. Strong track record of delivering end-to-end features across AWS, Azure, and GCP environments, with deep proficiency in React, SQL databases, and CI/CD pipelines. Passionate about building extensible platforms that enable teams to work seamlessly with AI agents while maintaining reliability and performance at scale.

Skills

Core Skills: React, HTMX, JavaScript, TypeScript, HTML5, CSS3, Redux, Next.js, Python, Node.js, REST APIs, GraphQL, Microservices, Distributed Systems, gRPC, AWS, Azure, GCP, Terraform, Ansible, Infrastructure as Code, Docker, Kubernetes, CI/CD, AI, ETL, Data Pipelines, Vector Databases, Feature Stores, LLM Integration, SQL, PostgreSQL, MongoDB, Observability, Monitoring, Prometheus, Grafana, ELK Stack, Datadog, Alerting, Logging, Jenkins, GitLab CI, GitHub Actions, Git, Linux, Agile, Scrum

Professional Experience

Google | Mountain View, CA Jan 2023 - Present

Senior Software Engineer, Infrastructure & AI Systems

Technologies: Python, React, TypeScript, GCP, Terraform, Kubernetes, PostgreSQL, Prometheus, AI, ETL

- Architected and deployed an AI-augmented workflow platform serving 500+ internal teams, integrating LLM capabilities with traditional backend systems using Python and React, resulting in 40% improvement in team productivity metrics.
- Designed comprehensive observability and monitoring infrastructure using Prometheus and Grafana, implementing real-time alerting for distributed data pipeline systems processing 10TB+ daily, reducing incident detection time from 2 hours to 5 minutes.
- Led development of feature store and vector database integration for AI systems, enabling seamless ETL pipelines that power LLM-based agents while maintaining 99.99% uptime across GCP infrastructure.
- Implemented infrastructure as code using Terraform and Ansible, automating deployment of microservices across multiple GCP regions, reducing infrastructure provisioning time by 70% and enabling self-service deployments for 15+ teams.

Microsoft | Seattle, WA Jun 2020 - Dec 2022

Software Engineer, Cloud Platform & Observability

Technologies: Python, React, Node.js, Azure, Terraform, Ansible, Docker, Kubernetes, SQL Server, Datadog, Monitoring

- Built end-to-end observability platform for distributed systems running on Azure, implementing custom monitoring dashboards in React that visualized metrics from 200+ microservices, enabling engineers to diagnose production issues 60% faster.
- Developed and maintained CI/CD pipelines using GitHub Actions and Azure DevOps, implementing infrastructure as code with Terraform and Ansible to manage 50+ cloud resources across development, staging, and production environments.

- Designed data pipeline architecture for AI model training, implementing ETL processes that ingested data from multiple sources into feature stores, supporting real-time inference for customer-facing AI features.
- Collaborated with cross-functional teams to gather requirements and deliver 12+ major features, demonstrating strong communication skills by presenting technical designs to both engineering and non-technical stakeholders.

Amazon | Seattle, WA Mar 2018 - May 2020

Software Engineer, Distributed Systems

Technologies: Python, JavaScript, React, AWS, Terraform, PostgreSQL, Kubernetes, Jenkins, Monitoring, Observability

- Engineered distributed database monitoring system handling 1M+ metrics per second, building React-based dashboards and implementing real-time alerting that reduced mean time to resolution by 50% for production incidents.
- Implemented infrastructure as code using Terraform to manage AWS resources, automating deployment of containerized microservices on Kubernetes and reducing manual configuration errors by 95%.
- Developed ETL pipelines for data warehouse ingestion, processing petabyte-scale datasets and optimizing SQL queries for complex analytical workloads supporting business intelligence teams.
- Mentored 5 junior engineers on distributed systems design patterns, CI/CD best practices, and production observability, fostering a culture of technical excellence and continuous improvement.

Stripe | San Francisco, CA Aug 2016 - Feb 2018

Full-Stack Software Engineer

Technologies: JavaScript, React, Node.js, Python, AWS, PostgreSQL, MongoDB, Docker, Jenkins, Git

- Developed full-stack payment processing features using React and Node.js, handling millions of transactions daily with 99.99% uptime, demonstrating deep understanding of distributed systems reliability.
- Built monitoring and alerting systems for payment infrastructure, implementing custom metrics and dashboards that provided real-time visibility into transaction flows and system health.
- Optimized PostgreSQL queries and database schema design for high-throughput systems, improving query performance by 40% and reducing database load during peak traffic periods.
- Contributed to CI/CD pipeline improvements using Jenkins, implementing automated testing and deployment strategies that reduced release cycle time from 4 hours to 30 minutes.

Uber | San Francisco, CA Jan 2015 - Jul 2016

Software Engineer, Backend Systems

Technologies: Python, Node.js, AWS, PostgreSQL, Microservices, Docker, Monitoring, Git

- Designed and implemented microservices for real-time ride matching system, processing 100K+ requests per second across distributed infrastructure while maintaining sub-100ms latency.
- Implemented comprehensive monitoring and observability solutions for backend systems, tracking performance metrics and implementing alerting rules that prevented 95% of potential customer-facing incidents.
- Optimized SQL queries and database indexing strategies for high-volume transactional systems, improving query response times by 35% and reducing infrastructure costs by 20%.
- Collaborated with product and data teams to understand requirements and deliver features on schedule, demonstrating excellent communication and ability to work effectively in fast-paced environments.

Certifications

- AWS Certified Solutions Architect - Professional
- AWS Certified DevOps Engineer - Professional
- Certified Kubernetes Administrator (CKA)
- HashiCorp Certified: Terraform Associate
- Google Cloud Professional Cloud Architect

- Microsoft Azure Solutions Architect Expert
- Datadog Certified Associate - Monitoring and Observability

Education

University of California, Berkeley

Master of Science in Computer Science 2014 **State University**

Bachelor of Science in Computer Science 2012